

Hadoop Yarn FifoScheduler Container 资源分配 Workflow

FifoScheduler

```
Public void handle(SchedulerEvent event) {
    ...
    case NODE_UPDATE: {
        NodeUpdateSchedulerEvent nodeUpdatedEvent =
            (NodeUpdateSchedulerEvent)event;
        nodeUpdate(nodeUpdatedEvent.getRMNode());
    }
    break;
    ...
}

private synchronized void nodeUpdate(RMNode rmNode) {
    ..
    // Processing the newly launched containers
    // process completed containers
    if (Resources.greaterThanOrEqualTo(resourceCalculator, clusterResource,
        node.getAvailableResource(), minimumAllocation)) {
        assignContainers(node);
    }
    metrics.setAvailableResourcesToQueue(
        Resources.subtract(clusterResource, usedResource));
}

private void assignContainers(FiCaSchedulerNode node) {
    for_each_app {
        synchronized(app) {
            for_each_priority {
                // ensure the application needs containers of this priority
                int assignedContainers =
                    assignContainersOnNode(node, application, priority);
            }
        }
    }
}

private int assignContainersOnNode(FiCaSchedulerNode node,
    FiCaSchedulerApp application, Priority priority) {
    // Data-local
    int nodeLocalContainers = assignNodeLocalContainers(node, application, priority);

    // Rack-local
    int rackLocalContainers = assignRackLocalContainers(node, application, priority);
}
```

```

// Off-switch
int offSwitchContainers = assignOffSwitchContainers(node, application, priority);
return (nodeLocalContainers + rackLocalContainers + offSwitchContainers);
}

private int assignNodeLocalContainers(FiCaSchedulerNode node,
    FiCaSchedulerApp application, Priority priority) {
    ...
    int assignableContainers =
        Math.min(
            getMaxAllocatableContainers(application, priority, node,
                NodeType.NODE_LOCAL),
            request.getNumContainers());
    assignedContainers =
        assignContainer(node, application, priority,
            assignableContainers, request, NodeType.NODE_LOCAL);
    return assignedContainers;
}

private int assignContainer(FiCaSchedulerNode node, FiCaSchedulerApp application,
    Priority priority, int assignableContainers,
    ResourceRequest request, NodeType type) {
    int availableContainers_mem =
        node.getAvailableResource().getMemory() / capability.getMemory();
    int availableContainers_cpu =
        node.getAvailableResource().getVirtualCores() / capability.getVirtualCores();

    int availableContainers =
        Math.min(availableContainers_mem, availableContainers_cpu);

    /*
    * @note GPU resource is not a must, so no modification about this upper bound
    */
    int assignedContainers = Math.min(assignableContainers, availableContainers);

    if (assignedContainers > 0) {
        for (int i=0; i < assignedContainers; ++i) {
            NodeId nodeId = node.getRMNode().getNodeID();
            ContainerId containerId = BuilderUtils.newContainerId(application
                .getApplicationAttemptId(), application.getNewContainerId());
            Token containerToken = null;
            /*
            * if this node has gpu resource and this request requests gpu, then assign it to this
            * container, otherwise, no gpu

```

```

        */
        if ( capability.getGPUCores() > node.getAvailableResource().getGPUCores())
        {
            capability=Resource.newInstance(capability);
            capability.setGPUCores(0); //set the require gpu as 0.
            capability.setGPUId(null);
        }

        /*
        * set the allocated gpu id, here we use only one gpu. 2013-12
        */
        else if(capability.getGPUCores()>0 && capability.getGPUCores()
            <= node.getAvailableResource().getGPUCores())
        {
            capability=Resource.newInstance(capability);
            int allocate_gpu_id=node.getAvailableResource().getGPUId().get(0);
            List<Integer> gpuld = new ArrayList<Integer>();
            gpuld.add(allocate_gpu_id);
            capability.setGPUId(gpuld);
        }
        containerToken = this.rmContext.getContainerTokenSecretManager()
            .createContainerToken(containerId, nodeId, application.getUser(),
                capability);
        if (containerToken == null) {
            return i; // Try again later.
        }

        // Create the container
        Container container =
            BuilderUtils.newContainer(containerId, nodeId, node.getRMNode()
                .getHttpAddress(), capability, priority, containerToken);

        // Inform the application
        RMContainer rmContainer =
            application.allocate(type, node, priority, request, container);

        // Inform the node
        node.allocateContainer(application.getApplicationId(), rmContainer);

        // Update usage for this container
        Resources.addTo(usedResource, capability);
    }
}

```